

Module 3

SEARCH IN COMPLEX ENVIRONMENTS

6 3.1 Local search algorithms

Local search algorithms are a class of optimization algorithms used in artificial intelligence (AI) to find solutions to optimization problems. These algorithms operate by iteratively exploring the solution space, gradually improving upon an initial solution until an optimal or satisfactory solution is found. Unlike global search algorithms that explore the entire solution space, local search algorithms focus on exploring a smaller subset of solutions, making them particularly useful for large or complex problem spaces where exhaustive search is impractical.

Some common local search algorithms used in AI include:

6
Hill Climbing: This algorithm iteratively improves a solution by making small incremental changes at each step that lead to a better solution. At each iteration, it evaluates neighboring solutions and moves to the neighboring solution with the highest improvement until a local maximum is reached.

4
Simulated Annealing: Inspired by the annealing process in metallurgy, this algorithm starts with a high probability of accepting worse solutions early in the search and gradually reduces this probability over time. This allows the algorithm to escape local optima and explore a wider range of solutions.

Genetic Algorithms (GA): Though often associated with global search, genetic algorithms can also be used as local search algorithms when applied to specific neighborhoods within the solution space. In each iteration, genetic algorithms generate new solutions by combining and mutating existing solutions, mimicking the process of natural selection.

39
Particle Swarm Optimization (PSO): PSO is inspired by the social behavior of bird flocks or fish schools. In PSO, a population of potential solutions (particles) moves through the solution space,

adjusting their positions based on their own best-known position and the global best-known position found by any particle in the swarm.

Local search algorithms are commonly applied to a variety of optimization problems in AI, including:

⁸⁶ **Traveling Salesman Problem (TSP):** Finding the shortest possible route that visits each city exactly once and returns to the origin city.

¹¹ **N-Queens Problem:** Placing N chess queens on an N×N chessboard such that no two queens threaten each other.

¹⁰⁴ **Graph Coloring Problem:** Assigning colors to vertices of a graph such that no two adjacent vertices have the same color with the fewest possible colors.

Job Scheduling: Assigning tasks to resources subject to constraints such as deadlines and resource availability.

¹¹ **Satisfiability Problem (SAT):** Determining whether a given Boolean formula can be satisfied by assigning Boolean values to its variables.

³¹ 3.2 Hill climbing Algorithm

Hill climbing is a simple local search algorithm that continually moves towards improving the current solution by making small incremental changes. It's a greedy algorithm that doesn't look ahead, and it terminates when it reaches a local maximum where no neighboring solution offers improvement.

⁹⁸ ⁴ let's illustrate the Hill-Climbing algorithm using the Traveling Salesman Problem (TSP) as an example. In the TSP, the goal is to find the shortest possible route that visits each city exactly once and returns to the starting city.

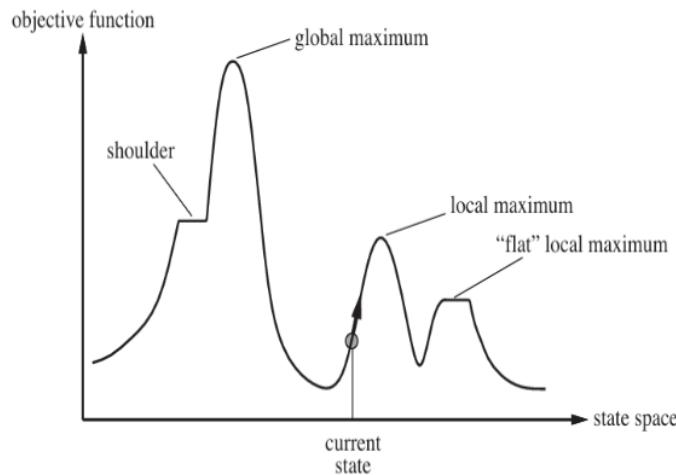


Figure 3.1 State space

98 Here's how the Hill-Climbing algorithm could be applied to the TSP:

Initialize: Start with a random initial solution. This could be any permutation of cities representing a possible route.

Evaluate: Calculate the total distance of the current route. This is the objective function we want to minimize.

4 **Generate neighbors:** Generate neighboring solutions by making small changes to the current solution. For example, swap the order of two cities in the route.

Select the best neighbor: Evaluate each neighbor and select the one that minimizes the objective function (i.e., has the shortest total distance).

Update: If the best neighbor improves upon the current solution, move to that neighbor and repeat the process. If no neighbor offers an improvement, terminate and return the current solution as the best-found solution.

14 **Repeat:** Continue steps 3-5 until a termination condition is met (e.g., a maximum number of iterations is reached, or no improvement is possible).

Let's discuss through a simple example:

165

Suppose we have 4 cities: A, B, C, and D. We start with the initial route ABCD:

Initial solution: ABCD

Total distance: Calculate the distance of the route (e.g., using Euclidean distance or any other distance metric).

We generate neighboring solutions by swapping the order of two cities:

Neighbor 1: BACD

Neighbor 2: ACBD

Neighbor 3: ABDC

Neighbor 4: ABCD (same as the current solution)

We evaluate the total distance for each neighbor and select the one with the shortest distance:

Neighbor 1 (BACD): Distance = 10

Neighbor 2 (ACBD): Distance = 12

Neighbor 3 (ABDC): Distance = 15

Neighbor 1 (BACD) has the shortest distance, so we move to this solution and repeat the process.

We continue this process until no further improvement is possible or until a termination condition is met.

If we get stuck in a local minimum where no neighbor offers an improvement, Hill-Climbing terminates without finding the optimal solution. To mitigate this, various enhancements such as random restarts or perturbation strategies can be applied.

Hill Climbing algorithm

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

```
current  $\leftarrow$  MAKE-NODE(problem.INITIAL-STATE)
loop do
    neighbor  $\leftarrow$  a highest-valued successor of current
    if neighbor.VALUE  $\leq$  current.VALUE then return current.STATE
    current  $\leftarrow$  neighbor
```

Hill climbing is a general optimization technique, and there are several variations of this algorithm designed to address different optimization scenarios. Some of the common types of hill climbing algorithms include:

Basic Hill Climbing: Also known as Simple Hill Climbing, this algorithm iteratively improves a solution by making small incremental changes at each step and moving to the neighboring solution with the highest improvement. It terminates when no better neighbor is found.

Steepest-Ascent Hill Climbing: In this variant, the algorithm evaluates all neighboring solutions and selects the one with the steepest ascent, i.e., the one that maximizes the improvement in the objective function. This ensures that the algorithm moves in the direction of the steepest gradient in the solution space. 85

Random-Restart Hill Climbing: This approach involves running the basic hill climbing algorithm multiple times from different initial random starting points. By restarting the search from different initial solutions, the algorithm increases the likelihood of finding the global optimum, especially in complex or rugged solution spaces.

Drawbacks

Local Optima: Hill climbing algorithms are prone to getting stuck in local optima, where the algorithm converges to a suboptimal solution without reaching the global optimum. Once the algorithm reaches a peak in the solution space, it cannot explore beyond that point, even if a better solution exists elsewhere. 91

Plateaus and Ridges: In rugged solution spaces characterized by plateaus (flat regions) or ridges (narrow paths), hill climbing algorithms struggle to navigate efficiently. When the algorithm encounters such regions, progress becomes slow or may halt altogether, leading to premature convergence.

Greedy Approach: Most hill climbing algorithms adopt a greedy strategy, always moving to the neighbor with the highest improvement in the objective function. While this strategy works well in some cases, it can lead to myopic decisions that prevent the algorithm from exploring promising regions that require sacrificing short-term gains for long-term benefits. 169 18

Limited Memory: Hill climbing algorithms typically have limited memory, focusing only on the current solution and its neighbors. This lack of global memory prevents the algorithm from retaining

information about previously visited solutions, hindering its ability to avoid revisiting the same solutions or exploring diverse regions of the solution space.³⁹

¹⁶ **3.3 Simulated Annealing**

Simulated Annealing is a probabilistic optimization algorithm inspired by the annealing process in metallurgy. It is used to find approximate solutions to optimization problems, particularly those involving complex and rugged solution spaces where traditional optimization algorithms may struggle.³² Simulated Annealing was introduced by Kirkpatrick, Gelatt, and Vecchi in 1983 and has since found applications in various fields, including combinatorial optimization, machine learning, and computational biology.⁴

Annealing Process Analogy: Simulated Annealing is inspired by the physical process of annealing in metallurgy, where a material is heated and then slowly cooled to reach a low-energy crystalline state. Similarly, in the optimization context, the algorithm starts with high exploration (high temperature)¹⁷⁹ and gradually decreases exploration (lowers temperature) to converge towards the optimal solution.

Objective Function: Like other optimization algorithms, Simulated Annealing aims to minimize or maximize an objective function that quantifies the quality of a solution. The objective function could represent, for example, the cost of a solution or its fitness in a given problem domain.³⁹

Neighbor Generation: At each iteration, Simulated Annealing generates a neighboring solution by making small perturbations to the current solution. The neighbors are generated probabilistically based on a neighborhood structure defined for the specific optimization problem.

Acceptance Criterion: Unlike greedy algorithms, Simulated Annealing accepts worse solutions with a certain probability during the early stages of the search.⁴ This probabilistic acceptance allows the algorithm to explore the solution space more effectively and escape local optima.¹²²

Temperature Schedule: The temperature parameter controls the degree of exploration in the algorithm. Initially high, the temperature gradually decreases over time according to a predefined schedule (annealing schedule). The cooling schedule determines how quickly the temperature decreases and influences the balance between exploration and exploitation.²²

Algorithm

```

function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to “temperature”

  current  $\leftarrow$  MAKE-NODE(problem.INITIAL-STATE)
  for t = 1 to  $\infty$  do
    T  $\leftarrow$  schedule(t)
    if T = 0 then return current
    next  $\leftarrow$  a randomly selected successor of current
     $\Delta E \leftarrow$  next.VALUE - current.VALUE
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{\Delta E/T}$ 

```

Workflow:

16

Initialization: Start with an initial solution and set an initial temperature.

102

Iterations: Repeat the following steps until a termination condition is met:

Generate a neighboring solution.

Evaluate the objective function of the new solution.

135

Determine whether to accept the new solution based on the acceptance criterion.

Update the temperature according to the cooling schedule.

10

Termination: Terminate the algorithm when a stopping criterion is satisfied (e.g., reaching a maximum number of iterations, achieving a desired solution quality).

Acceptance Criterion (Metropolis criterion):

22

If the new solution is better than the current solution, accept it.

If the new solution is worse, accept it with a probability

Applications:

Combinatorial Optimization: Traveling Salesman Problem, Graph Coloring, Job Scheduling.

Machine Learning: Training neural networks, optimizing hyperparameters.

Physical Systems: Modeling complex systems in physics and chemistry.

Robotics: Path planning for robots in complex environments.

Simulated Annealing example: The Traveling Salesman Problem (TSP). In this example, we'll consider a small set of cities and aim to find the shortest possible route that visits each city exactly once and returns to the starting city.⁸³

Example: Traveling Salesman Problem with Simulated Annealing

Problem Setup:

Let's consider a scenario with 5 cities (A, B, C, D, E) and their respective coordinates:⁴

A: (0, 0)¹⁸⁷

B: (1, 2)

C: (3, 1)

D: (5, 2)

E: (4, 0)

Objective Function:

The objective function is the total distance of the route. We'll use the Euclidean distance formula to calculate distances between cities.¹³⁸⁹

Simulated Annealing Algorithm:

Initialization:

Start with a random initial solution (random permutation of cities).¹¹⁹

Set the initial temperature

T and cooling rate.

Iterations:

Repeat until the stopping criterion is met:

Generate a neighboring solution by randomly swapping two cities in the current route.²²

Calculate the change in objective

If the new solution is better (lower distance), accept it.

If the new solution is worse, accept it with a probability determined by the acceptance criterion.

Update the temperature according to the cooling schedule.

Termination:

Terminate the algorithm when a stopping criterion is met (e.g., reaching a maximum number of iterations, convergence to a desired solution quality).¹⁰

Implementation:

Let's simulate the annealing process step by step:

Initialization:

Start with a random initial solution, e.g., ABCDE.

Set initial temperature

T and cooling rate.

Iterations:

22

Generate a neighboring solution by swapping two cities, e.g., ACBDE.

Calculate the change in distance

80

If the new solution is better, accept it.

If the new solution is worse, accept it with a probability

Update the temperature and repeat the process.

Termination:

152

Stop when the temperature reaches a minimum value or after a fixed number of iteration.

59

3.4 Local Beam Search

47

Local Beam Search is a variation of the beam search algorithm used for solving optimization problems. In this algorithm, a fixed number of candidate solutions, known as the beam width, are explored simultaneously at each iteration. It is particularly useful for problems where the solution space is large and the goal is to find a satisfactory solution quickly.

Steps:

50

Initialization: Begin with generating k initial candidate solutions randomly or using a heuristic method.

Generation of Neighbors: For each candidate solution, generate neighboring solutions by applying local moves or perturbations. These neighbors can be generated by making small modifications to the current solutions, such as swapping elements or applying operators specific to the problem domain.

24

Evaluation: Evaluate the quality of each generated neighbor using an objective function or evaluation metric. This function measures how well a solution performs with respect to the problem's goals.

Selection: Select the top k solutions from the pool of generated neighbors based on their evaluation scores. These solutions become the candidates for the next iteration.

Termination Criteria: Repeat the process iteratively until a termination criterion is met. This criterion can be a maximum number of iterations, reaching a desired solution quality, or any other condition specific to the problem.

Algorithm Steps:

Initialization:

Generate k initial candidate solutions randomly or using a heuristic method.

Iterations:

Generate neighboring solutions for each candidate solution.

Evaluate the quality of each neighbor using an objective function.

Select the top k neighbors based on their evaluation scores.

Termination:

Terminate the algorithm when a termination criterion is met.

Advantages of Local Beam Search:

Diverse Exploration: Since multiple candidate solutions are maintained simultaneously, local beam search can explore a broader region of the solution space compared to algorithms that focus on a single solution.

Efficiency: Local beam search can quickly converge to good solutions, especially in problems where the solution space is vast and exhaustive exploration is impractical.

Parallelization: The algorithm naturally lends itself to parallelization, as each candidate solution can be processed independently in parallel.

Disadvantages of Local Beam Search:

24

Limited Exploration: Local beam search may get stuck in local optima, especially if the initial set of candidate solutions does not cover a diverse range of the solution space.

77

Beam Width Selection: The performance of the algorithm can be sensitive to the choice of beam width.

77

k. A small k may lead to premature convergence, while a large k may increase computational overhead.

102

Memory Requirements: Maintaining a large number of candidate solutions in memory can be resource-intensive, especially for problems with high-dimensional solution spaces.

Applications:

Local beam search is commonly used in various optimization problems, including:

113

Constraint Satisfaction Problems (CSPs): Assigning values to variables subject to constraints.

Job Scheduling: Allocating tasks to resources while optimizing certain criteria.

Route Optimization: Finding the shortest or most efficient routes in transportation and logistics.

24

Overall, local beam search is a powerful optimization algorithm suitable for a wide range of problems, particularly those with large solution spaces where exploration is crucial for finding satisfactory solutions.

3.5 Evolutionary Algorithms

54

Evolutionary Algorithms (EAs) are a class of optimization algorithms inspired by the principles of biological evolution. These algorithms mimic the process of natural selection to iteratively search for solutions to optimization problems. Evolutionary algorithms are widely used in various fields, including engineering, computer science, economics, and biology, to solve complex optimization problems.

24

Population: Evolutionary algorithms maintain a population of candidate solutions (individuals or chromosomes) representing potential solutions to the optimization problem.

4

Fitness Function: A fitness function evaluates the quality of each candidate solution in the population. It quantifies how well a solution performs with respect to the problem's objectives.

41

Selection: Candidates are selected from the population for reproduction based on their fitness scores. Solutions with higher fitness scores are more likely to be selected, mimicking the principle of "survival of the fittest."

Reproduction: Selected candidates undergo genetic operations such as crossover and mutation to produce offspring. These operations combine genetic material from parent solutions to create new solutions with potentially improved characteristics.

Replacement: Offspring replace existing solutions in the population based on criteria such as elitism or generational turnover. This maintains the population size and diversity over successive generations.

10

Termination Criteria: Evolutionary algorithms terminate when a stopping criterion is met, such as reaching a maximum number of generations, achieving a desired solution quality, or exceeding a computational budget.

24

Common Variants of Evolutionary Algorithms:

4

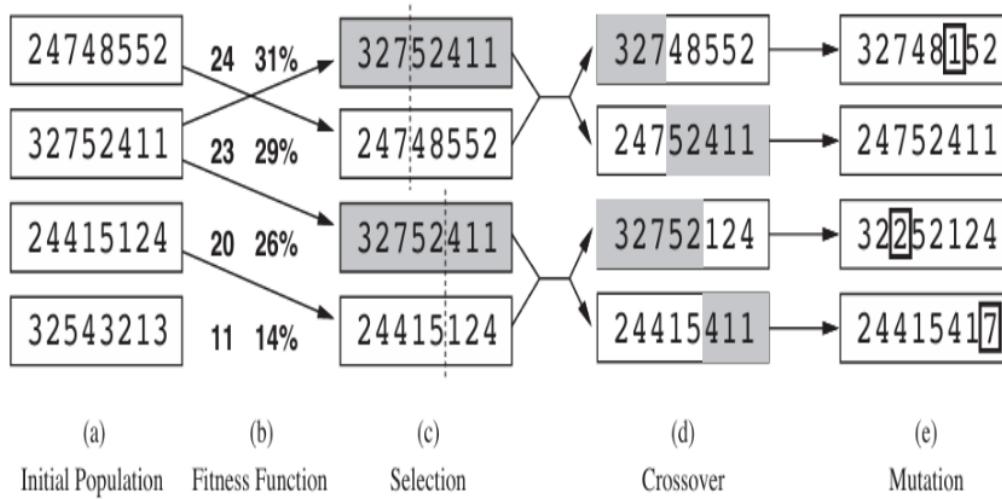
Genetic Algorithms (GA): Genetic algorithms use techniques inspired by genetics and natural selection, such as crossover, mutation, and selection operators, to evolve a population of solutions over multiple generations.

Evolutionary Strategies (ES): Evolutionary strategies focus on directly optimizing the parameters of a solution using mutation and selection, often in high-dimensional continuous search spaces.

39

Genetic Programming (GP): Genetic programming evolves programs or models represented as trees rather than fixed-length chromosomes, making it suitable for symbolic regression, function optimization, and automatic program synthesis.

Algorithm



24

Differential Evolution (DE): Differential evolution optimizes a population of candidate solutions by iteratively combining the differences between randomly selected individuals.

10

211

Particle Swarm Optimization (PSO): Although not strictly an evolutionary algorithm, PSO is inspired by the social behavior of bird flocks or fish schools. It maintains a population of particles that move through the solution space, adjusting their positions based on their own and their neighbors' best-known positions.

56

Applications:

Evolutionary algorithms find applications in various optimization problems, including:

10

Function Optimization: Finding the global minimum or maximum of a mathematical function.

Machine Learning: Tuning hyperparameters of machine learning algorithms and neural networks.

Engineering Design: Optimizing designs in mechanical, electrical, and aerospace engineering.

Robotics: Path planning, motion control, and robot learning.

Bioinformatics: Protein folding, sequence alignment, and metabolic pathway optimization.

Advantages of Evolutionary Algorithms:

Global Search: Evolutionary algorithms are well-suited for exploring large and complex solution spaces, making them effective for finding global optima.

Versatility: They can handle various types of optimization problems, including continuous, discrete, combinatorial, and multi-objective problems.

Parallelization: Evolutionary algorithms can be parallelized easily, enabling efficient exploration of the solution space using parallel and distributed computing resources.

Disadvantages of Evolutionary Algorithms:

Computational Cost: Evolutionary algorithms can be computationally expensive, especially for high-dimensional or multimodal optimization problems.

Parameter Tuning: The performance of evolutionary algorithms is sensitive to parameter settings such as population size, mutation rate, and selection strategy, requiring careful tuning.

Premature Convergence: Evolutionary algorithms may converge prematurely to suboptimal solutions if the population lacks diversity or the exploration-exploitation balance is not maintained.

Despite their limitations, evolutionary algorithms remain powerful tools for solving complex optimization problems where traditional methods may struggle to find satisfactory solutions.

3.6 Optimal decision-making in games

Optimal decision-making in games is a fundamental concept in game theory and artificial intelligence, particularly in the context of competitive games where players aim to maximize their chances of winning. Various strategies and algorithms are employed to make optimal decisions in games, depending on the complexity of the game, the available information, and the goals of the players.

The Minimax algorithm is a decision-making algorithm commonly used in two-player zero-sum games, such as chess, checkers, and tic-tac-toe. It aims to find the optimal move for a player by exploring the entire game tree, considering all possible future moves and their outcomes. The algorithm assumes that both players play optimally, with one player attempting to maximize their score (the maximizing player), and the other player attempting to minimize the score (the minimizing player).

Game Tree: The Minimax algorithm operates on a game tree that represents all possible moves and their outcomes. ² The root of the tree corresponds to the current game state, and each node represents a possible state resulting from a player's move.

Maximizing Player: ⁴³ The player whose turn it is to move seeks to maximize their score (utility). This player aims to choose the move that leads to the highest possible utility among all possible moves.

Minimizing Player: The opponent player aims to minimize the maximizing player's score. They choose moves that lead to the lowest possible utility for the maximizing player.

Recursive Search: Minimax performs a depth-first search of the game tree, evaluating each node recursively. ¹ At each level of the tree, the algorithm alternates between maximizing ³³ and minimizing ⁴⁰ players until it reaches a terminal state (end of the game).

Utility Function: The utility function assigns a numerical value to each terminal state, representing the outcome of the game (e.g., win, loss, draw). The goal is to maximize the utility for the maximizing player and minimize it for the minimizing player.

Algorithm Steps:

Recursive Search:

¹⁸ Start at the root node representing the current game state.

Recursively explore the game tree, alternating between maximizing and minimizing players.

At each level of the tree:

⁶⁷ If it's the maximizing player's turn, choose the child node with the highest utility.

If it's the minimizing player's turn, choose the child node with the lowest utility.

⁴³ Continue until reaching a terminal state (end of the game).

Terminal State Evaluation:

¹ When the algorithm reaches a terminal state (e.g., win, loss, draw), evaluate the utility of that state using the utility function.

Backpropagation:

²⁹ Propagate the utility value of terminal states back up the tree to update the utility values of parent nodes.

At each parent node, update the utility value based on the maximum or minimum utility of its child nodes, depending on whether it represents a maximizing or minimizing player.

Decision Making:

32

Once the entire tree is evaluated, choose the move that leads to the highest utility for the maximizing player from the root node.

Consider a simple example of the Tic-Tac-Toe game. In this game:

The maximizing player aims to win the game or achieve a draw.

The minimizing player aims to prevent the maximizing player from winning.

The utility function assigns values (+1 for win, 0 for draw, -1 for loss) to terminal states.

The Minimax algorithm explores the entire game tree of Tic-Tac-Toe, evaluating each possible move and its outcome until reaching a terminal state. Based on the computed utility values, it selects the move that maximizes the maximizing player's chances of winning or achieving a draw, assuming optimal play by both players.

3.7 Optimal decisions in multiplayer games

In multiplayer games, optimal decision-making becomes more complex compared to two-player games due to the increased number of opponents and interactions among players. While the Minimax algorithm is a fundamental concept in decision-making for two-player zero-sum games, its direct application to multiplayer games is challenging. Instead, various strategies and algorithms are used to make optimal decisions in multiplayer games, depending on factors such as game dynamics, player interactions, and strategic objectives.

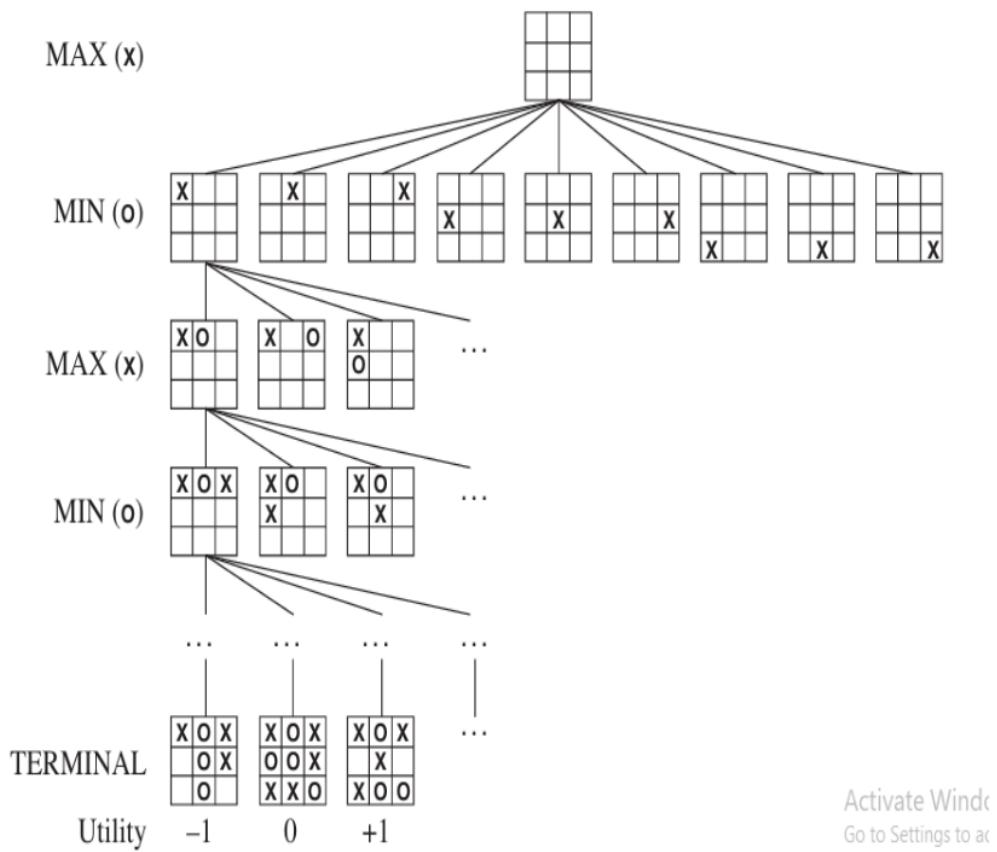


Figure : Tic-Tac-Toe

40 **3.8 Alpha-Beta pruning** is a powerful optimization technique used in decision trees, particularly in game trees, to reduce the number of nodes that need to be evaluated during the Minimax algorithm. 40

45 It aims to decrease the number of nodes visited in the search tree by eliminating branches that are guaranteed not to influence the final decision. This optimization significantly improves the efficiency of the Minimax algorithm, making it feasible to search deeper into the game tree.

51 **Alpha and Beta Values:**

161 Alpha represents the best (highest) value found so far by any maximizing player along the path from the root to the current node. 15

Beta represents the best (lowest) value found so far by any minimizing player along the path from the root to the current node.

Pruning Condition:

⁹⁶ When a node's value exceeds the beta value of its parent (for a minimizing player) or falls below the alpha value of its parent (for a maximizing player), the search can be pruned. This is because the parent node will never choose this node, as there exists a better alternative elsewhere.

Recursive Algorithm:

The alpha-beta pruning algorithm is typically implemented recursively, evaluating nodes in a depth-first manner while updating ⁴³ alpha and beta values at each level of the tree.

Efficiency:

⁴⁵ Alpha-beta pruning can drastically reduce the number of nodes evaluated in the search tree, especially in games with high branching factors or deep search depths. By eliminating irrelevant branches early in the search process, the algorithm can focus on promising paths, leading to significant computational savings.

Algorithm Steps:

Initialization:

Start with initial alpha and beta values representing negative and positive infinity, respectively.

Recursion: Recursively traverse the game tree, evaluating nodes and updating alpha and beta values at each level.

⁵¹ **Pruning:** If, at any point during the traversal, a node's value exceeds its parent's beta value (for a minimizing player) or falls below its parent's alpha value (for a maximizing player), prune the subtree rooted at that node.

Base Case: When reaching leaf nodes (terminal states or maximum search depth), evaluate the node's value directly.

Return: Return the best value found (alpha for maximizing player, beta for minimizing player) up the tree.

¹⁹² **Example:** Let's consider a simple game tree with two players, Max and Min, and a search depth of 3 levels. Here's how the alpha-beta pruning algorithm works:

Start at the root node with initial alpha and beta values.

Recursively explore child nodes, updating alpha and beta values.

Prune subtrees when possible based on the pruning condition.

Continue until reaching leaf nodes or the maximum search depth.

Return the best value found (alpha or beta) up the tree.

By pruning irrelevant branches early in the search process, alpha-beta pruning can significantly reduce the number of nodes evaluated, making it possible to search deeper into the game tree and find more optimal solutions efficiently.

In summary, alpha-beta pruning is a fundamental optimization technique used in decision trees, particularly in game trees, to improve the efficiency of the Minimax algorithm. It allows for more effective exploration of the solution space by eliminating irrelevant branches, leading to faster and more accurate decision-making in games and other search problems.

3.9 Move Ordering:

Alpha-Beta pruning reduces the number of nodes evaluated in the search tree by eliminating branches that are guaranteed not to affect the final decision. By ordering the moves properly, the algorithm can potentially prune more branches, leading to faster search times.

Pruning Efficiency: Proper move ordering increases the likelihood of encountering favorable branches early in the search, allowing Alpha-Beta pruning to prune more nodes efficiently. This reduces the effective branching factor and speeds up the search process.

Deeper Cutoffs: Effective move ordering increases the chances of reaching deeper cutoffs in the search tree, enabling the algorithm to explore more promising branches and improve the quality of the final decision.

Improved Pruning Heuristics: Move ordering can be combined with other pruning heuristics, such as history heuristic or killer heuristic, to guide the search towards more promising moves and enhance pruning efficiency further.

Techniques for Move Ordering:

Captures and Promotions: In games like chess, moves that involve captures or promotions often have a significant impact on the game's outcome. Prioritizing these moves in the search can lead to more effective pruning.

Good and Bad Captures: Prioritize captures that result in material gain and avoid captures that result in material loss. This heuristic can guide the search towards moves that are likely to improve the player's position.

Killer Heuristic: Keep track of moves that cause cutoffs at specific depths in previous iterations of the search. These "killer moves" are likely to be strong moves in the current position and can be considered early in the search.

History Heuristic: Maintain a history table that records the frequency of successful moves in previous searches. Moves with a high success rate are likely to be good candidates for exploration and can be prioritized.

Static Evaluation Function: Use a static evaluation function to estimate the quality of moves at a shallow search depth. The evaluation function can guide move ordering by prioritizing moves that lead to positions with higher scores.

Principal Variation (PV) Move: The principal variation represents the best line of play identified in the search. Prioritize moves along the principal variation to explore promising branches first.

Implementation:

Ordering Phase: Before starting the search, generate and order the moves based on the chosen move ordering heuristic.

Search Phase: During the search, explore the moves in the order determined by the move ordering heuristic. This ensures that more promising moves are explored first, potentially leading to deeper cutoffs and more efficient pruning.

Iterative Refinement: Continuously refine the move ordering heuristic based on the results of previous searches. Adjust the heuristic parameters and update the move ordering strategies to improve search performance over time.

³⁸ **3.10 Monte Carlo Tree Search (MCTS)**

Monte Carlo Tree Search (MCTS) is a powerful algorithm used in artificial intelligence for decision-making in games and other domains with large state spaces and complex decision trees. It has gained

prominence in recent years for its success in games like Go, Poker, and various video games. MCTS combines random sampling with tree search to efficiently explore the game tree and make decisions. Here's how MCTS works:

90 1. Selection Phase:

Starting from the root node of the search tree, the algorithm selects child nodes iteratively until it reaches a leaf node. The selection process balances between exploring promising branches and exploiting existing knowledge.

38 2. Expansion Phase:

Once a leaf node is reached, the algorithm expands the tree by adding one or more child nodes representing possible future game states. The expansion phase increases the breadth and depth of the search tree, allowing the algorithm to explore new possibilities.

1. Simulation (Rollout) Phase:

18 From the newly added nodes, the algorithm performs a simulation or rollout until reaching a terminal state (end of the game). During the simulation, actions are chosen randomly or according to a simple policy to explore the possible outcomes of the game.

1. Backpropagation Phase:

After completing the simulation, the algorithm updates the statistics of visited nodes and propagates the outcome back up the tree. The backpropagation phase incrementally updates the win/loss statistics and visit counts of nodes along the path from the leaf node to the root node.

Node Selection: MCTS uses selection strategies to decide which child nodes to explore further. Popular selection strategies include the Upper Confidence Bound (UCB) algorithm, which balances between exploration (visiting less explored nodes) and exploitation (visiting nodes with high estimated value).

Expansion Policy: The expansion policy determines how new nodes are added to the tree. Typically, nodes are expanded based on legal moves or actions available in the game.

Rollout Policy: The rollout policy dictates how actions are chosen during the simulation phase. While the rollout policy can be simple (e.g., random moves), more sophisticated policies may use heuristics or domain-specific knowledge.

Backpropagation Update Rule: The backpropagation phase updates the statistics of visited nodes based on the outcome of the simulation. The update rule depends on the specific application and can vary, but commonly used methods include averaging or incrementing win/loss counts.

Advantages of MCTS:

Domain Independence: MCTS is applicable to a wide range of domains, including board games, video games, and combinatorial optimization problems.

Asymptotic Optimality: As the search progresses, MCTS converges to the optimal solution under certain conditions, making it suitable for finding near-optimal solutions in complex problems.

Adaptive Search: MCTS dynamically allocates search effort to promising regions of the search space, allowing it to focus on relevant areas and ignore irrelevant ones.

Applications of MCTS:

Games: MCTS has been successfully applied to games like Go, Chess, Shogi, Poker, and various video games.

Planning and Optimization: MCTS can be used for planning and optimization problems with large state spaces, such as route planning, scheduling, and resource allocation.

Robotics: MCTS is used in robotics for decision-making in uncertain environments, path planning, and autonomous navigation.

3.11 The Kalman Filter

The Kalman Filter is a mathematical algorithm used for state estimation in systems with uncertain dynamic models and noisy measurements. While it's not exclusively an AI algorithm, it's widely used in AI and robotics for various tasks such as tracking, localization, sensor fusion, and control.

Basic Concept:

The Kalman Filter estimates the state of a linear dynamic system over time based on a series of noisy measurements. It combines predictions from a dynamic model of the system with measurements from sensors to produce a more accurate estimate of the true state.

Components of the Kalman Filter:

186

State Transition Model (Prediction): Represents the evolution of the system over time. Typically described as a linear dynamical system: It is the measurement noise.

9

State Estimate: Represents the estimated state of the system at each time step.

53

Computed using a combination of the predicted state and the measurement update.

Error Covariance: Represents the uncertainty in the state estimate.

Updated based on the accuracy of the predictions and the reliability of the measurements.